

GitLab CI: Run jobs sequentially, in parallel or build a custom pipeline

Ivan Nemytchenko (<https://twitter.com/inemation>) Jul 29, 2016

Let's assume that you don't know anything about what Continuous Integration is and why it's needed. Or, you just forgot. Anyway, we're starting from scratch here.

Imagine that you work on a project, where all the code consists of two text files. Moreover, it is super-critical that the concatenation of these two files contains the phrase "Hello world."

If there's no such phrase, the whole development team stays without a salary for a month. Yeah, it is that serious!

The most responsible developer wrote a small script to run every time we are about to send our code to customers. The code is pretty sophisticated:

```
cat file1.txt file2.txt | grep -q "Hello world"
```

The problem is that there are ten developers in the team, and, you know, human factors can hit hard.

A week ago, a new guy forgot to run the script and three clients got broken builds. So you decided to solve the problem once and for all. Luckily, your code is already on GitLab, and you remember that there is a built-in CI system (<https://about.gitlab.com/gitlab-ci/>). Moreover, you heard at a conference that people use CI to run tests...

Run our first test inside CI

After a couple minutes to find and read the docs, it seems like all we need is these two lines of code in a file called `.gitlab-ci.yml`:

```
test:
  script: cat file1.txt file2.txt | grep -q 'Hello world'
```

Committing it, and hooray! Our build is successful:

 Build #2346110 for commit f8a26206 from master by @inem about an hour ago

```
gitlab-ci-multi-runner 1.3.2 (0323456)
Using Docker executor with image ruby:2.1 ...
Pulling docker image ruby:2.1 ...
Running on runner-8a2f473d-project-1398078-concurrent-0 via runner-8a2f473d-machine-1468420047-2374f4cc-digital-ocean-4gb...
Cloning repository...
Cloning into '/builds/inem/ci'...
Checking out f8a26206 as master...
$ cat file1.txt file2.txt | grep -q "Hello world"

Build succeeded
```

Let's change "world" to "Africa" in the second file and check what happens:

 Build #2346623 for commit b978b9f6 from master by @inem about an hour ago

```
gitlab-ci-multi-runner 1.3.2 (0323456)
Using Docker executor with image ruby:2.1 ...
Pulling docker image ruby:2.1 ...
Running on runner-30dcea4b-project-1398078-concurrent-0 via runner-30dcea4b-machine-1468421193-15f1e5c5-digital-ocean-4gb...
Cloning repository...
Cloning into '/builds/inem/ci'...
Checking out b978b9f6 as master...
$ cat file1.txt file2.txt | grep -q "Hello world"

ERROR: Build failed: exit code 1
```

The build fails as expected!

Okay, we now have automated tests here! GitLab CI will run our test script every time we push new code to the repository.

Make results of builds downloadable

The next business requirement is to package the code before sending it to our customers. Let's automate that as well!

All we need to do is define another job for CI. Let's name the job "package":

```
test:
  script: cat file1.txt file2.txt | grep -q 'Hello world'

package:
  script: cat file1.txt file2.txt | gzip > package.gz
```

We have two tabs now:

✓ test ✓ package

```
gitlab-ci-multi-runner 1.3.2 (0323456)
Using Docker executor with image ruby:2.1 ...
Pulling docker image ruby:2.1 ...
Running on runner-8a2f473d-project-1447361-concurrent-0 via runner-8a2f473d-machine-1469549805-b2f018ac-digital-ocean-4gb...
Cloning repository...
Cloning into '/builds/inem/ci2'...
Checking out be833f45 as master...
$ cat file1.txt file2.txt | grep -q 'Hello world'

Build succeeded
```

However, we forgot to specify that the new file is a build *artifact*, so that it could be downloaded. We can fix it by adding an `artifacts` section:

```
test:
  script: cat file1.txt file2.txt | grep -q 'Hello world'

package:
  script: cat file1.txt file2.txt | gzip > packaged.gz
  artifacts:
    paths:
      - packaged.gz
```

Checking... It is there:

✓ passed Build #2630934 for commit 0e99231d from master by @inem 2 minutes ago

✓ test ✓ package

```
gitlab-ci-multi-runner 1.3.2 (0323456)
Using Docker executor with image ruby:2.1 ...
Pulling docker image ruby:2.1 ...
Running on runner-30dcea4b-project-1447361-concurrent-0 via runner-30dcea4b-machine-1469550348-1f5e7833-digital-ocean-4gb...
Cloning repository...
Cloning into '/builds/inem/ci2'...
Checking out 0e99231d as master...
$ cat file1.txt file2.txt | gzip > packaged.gz
Uploading artifacts...
packaged.gz: found 1 matching files
Uploading artifacts to coordinator... ok      id=2630934 responseStatus=201 Created token=7PE1aQwt

Build succeeded
```

Build artifacts

Download Browse

Build details Retry

Duration: 49 seconds

Finished: about a minute ago

Runner: #21099

Raw Erase

Commit title

Update .gitlab-ci.yml

Perfect! However, we have a problem to fix: the jobs are running in parallel, but we do not want to package our application if our tests fail.

Run jobs sequentially

We only want to run the 'package' job if the tests are successful. Let's define the order by specifying stages :

```
stages:
  - test
  - package

test:
  stage: test
  script: cat file1.txt file2.txt | grep -q 'Hello world'

package:
  stage: package
  script: cat file1.txt file2.txt | gzip > packaged.gz
  artifacts:
    paths:
      - packaged.gz
```

That should be good!

Also, we forgot to mention, that compilation (which is represented by concatenation in our case) takes a while, so we don't want to run it twice. Let's define a separate step for it:

```
stages:
  - compile
  - test
  - package

compile:
  stage: compile
  script: cat file1.txt file2.txt > compiled.txt
  artifacts:
    paths:
      - compiled.txt

test:
  stage: test
  script: cat compiled.txt | grep -q 'Hello world'

package:
  stage: package
  script: cat compiled.txt | gzip > packaged.gz
  artifacts:
    paths:
      - packaged.gz
```

Let's take a look at our artifacts:



Download 'compile' artifacts



Download 'package' artifacts

Hmm, we do not need that "compile" file to be downloadable. Let's make our temporary artifacts expire by setting `expire_in` to '20 minutes':

```
compile:
  stage: compile
  script: cat file1.txt file2.txt > compiled.txt
  artifacts:
    paths:
      - compiled.txt
    expire_in: 20 minutes
```

Now our config looks pretty impressive:

- We have three sequential stages to compile, test, and package our application.
- We are passing the compiled app to the next stages so that there's no need to run compilation twice (so it will run faster).
- We are storing a packaged version of our app in build artifacts for further usage.

Learning which Docker image to use

So far so good. However, it appears our builds are still slow. Let's take a look at the logs.

```

gitlab-ci-multi-runner 1.3.2 (0323456)
Using Docker executor with image ruby:2.1 ...
Pulling docker image ruby:2.1 ...
Running on runner-30dcea4b-project-1398078-concurrent-0 via runner-30dcea4b-machine-1469178408-7ba044d5-digital-ocean-4gb...
Cloning repository...
Cloning into '/builds/inem/ci'...
Checking out bdc26c45 as master...
$ cat file1.txt file2.txt > compiled.txt
Uploading artifacts...
compiled.txt: found 1 matching files
Uploading artifacts to coordinator... ok      id=2545753 responseStatus=201 Created token=39RXWAPw

Build succeeded
  
```

Wait, what is this? Ruby 2.1?

Why do we need Ruby at all? Oh, GitLab.com uses Docker images to run our builds (/2016/04/05/shared-runners/), and by default (/gitlab-com/settings/#shared-runners) it uses the `ruby:2.1` (https://hub.docker.com/_/ruby/) image. For sure, this image contains many packages we don't need. After a minute of googling, we figure out that there's an image called `alpine` (https://hub.docker.com/_/alpine/) which is an almost blank Linux image.

OK, let's explicitly specify that we want to use this image by adding `image: alpine` to `.gitlab-ci.yml`. Now we're talking! We shaved almost 3 minutes off:

Status	Commit	Compile	Test	Package	
✓ passed	#3795973 master ↗ 628c2511 latest Update .gitlab-ci.yml	✓	✓	✓	⌚ 00:35 📅 about a minute ago
✓ passed	#3792115 master ↗ 4b5604cc Update .gitlab-ci.yml	✓	✓	✓	⌚ 03:27 📅 about 14 hours ago

It looks like there's (https://hub.docker.com/_/mysql/) a lot of (https://hub.docker.com/_/python/) public images (https://hub.docker.com/_/java/) around (https://hub.docker.com/_/php/). So we can just grab one for our technology stack. It makes sense to specify an image which contains no extra software because it minimizes download time.

Dealing with complex scenarios

So far so good. However, let's suppose we have a new client who wants us to package our app into `.iso` image instead of `.gz`. Since CI does the whole work, we can just add one more job to it. ISO images can be created using the `mkisofs` (http://linuxcommand.org/man_pages/mkisofs8.html) command. Here's how our config should look:

```

image: alpine

stages:
  - compile
  - test
  - package

# ... "compile" and "test" jobs are skipped here for the sake of compactness

pack-gz:
  stage: package
  script: cat compiled.txt | gzip > packaged.gz
  artifacts:
    paths:
      - packaged.gz

pack-iso:
  stage: package
  script:
    - mkisofs -o ./packaged.iso ./compiled.txt
  artifacts:
    paths:
      - packaged.iso

```

Note that job names shouldn't necessarily be the same. In fact if they were the same, it wouldn't be possible to make the jobs run in parallel inside the same stage. Hence, think of same names of jobs & stages as coincidence.

Anyhow, the build is failing:

✓ compile
✓ test
✓ pack-gz
✗ pack-iso

```

gitlab-ci-multi-runner 1.3.2 (0323456)
Using Docker executor with image alpine ...
Pulling docker image alpine ...
Running on runner-30dcea4b-project-1447361-concurrent-0 via runner-30dcea4b-machine-1469599207-54a20bb2-digital-ocean-4gb...
Cloning repository...
Cloning into '/builds/inem/ci2'...
Checking out 7d771ca2 as master...
Downloading artifacts for compile (2642195)...
Downloading artifacts from coordinator... ok      id=2642195 responseStatus=200 OK token=7PE1aQwt
$ mkisofs -o ./packaged.iso ./compiled.txt
/bin/sh: eval: line 40: mkisofs: not found

ERROR: Build failed: exit code 127

```

The problem is that `mkisofs` is not included in the `alpine` image, so we need to install it first.

Dealing with missing software/packages

According to the Alpine Linux website (<https://pkgs.alpinelinux.org/contents?file=mkisofs&path=&name=&branch=&repo=&arch=x86>) `mkisofs` is a part of the `xorriso` and `cdrkit` packages. These are the magic commands that we need to run to install a package:

```
echo "ipv6" >> /etc/modules # enable networking
apk update                 # update packages list
apk add xorriso            # install package
```

For CI, these are just like any other commands. The full list of commands we need to pass to `script` section should look like this:

```
script:
- echo "ipv6" >> /etc/modules
- apk update
- apk add xorriso
- mkisofs -o ./packaged.iso ./compiled.txt
```

However, to make it semantically correct, let's put commands related to package installation in `before_script`. Note that if you use `before_script` at the top level of a configuration, then the commands will run before all jobs. In our case, we just want it to run before one specific job.

Our final version of `.gitlab-ci.yml`:


```

image: alpine

stages:
  - compile
  - test
  - package

compile:
  stage: compile
  script: cat file1.txt file2.txt > compiled.txt
  artifacts:
    paths:
      - compiled.txt
    expire_in: 20 minutes


test:
  stage: test
  script: cat compiled.txt | grep -q 'Hello world'

pack-gz:
  stage: package
  script: cat compiled.txt | gzip > packaged.gz
  artifacts:
    paths:
      - packaged.gz

pack-iso:
  stage: package
  before_script:
    - echo "ipv6" >> /etc/modules
    - apk update
    - apk add xorriso
  script:
    - mkisofs -o ./packaged.iso ./compiled.txt
  artifacts:
    paths:
      - packaged.iso

```

Wow, it looks like we have just created a pipeline! We have three sequential stages, but jobs `pack-gz` and `pack-iso`, inside the `package` stage, are running in parallel:



Summary

There's much more to cover but let's stop here for now. I hope you liked this short story. All examples were made intentionally trivial so that you could learn the concepts of GitLab CI without being distracted by an unfamiliar technology stack. Let's wrap up what we have learned:

1. To delegate some work to GitLab CI you should define one or more jobs (<http://docs.gitlab.com/ce/ci/yaml/README.html#jobs>) in `.gitlab-ci.yml`.
2. Jobs should have names and it's your responsibility to come up with good ones.
3. Every job contains a set of rules & instructions for GitLab CI, defined by special keywords.
4. Jobs can run sequentially, in parallel, or you can define a custom pipeline.
5. You can pass files between jobs and store them in build artifacts so that they can be downloaded from the interface.

Below is the last section containing a more formal description of terms and keywords we used, as well as links to the detailed description of GitLab CI functionality.

Keywords description & links to the documentation

Keyword/term	Description
<code>.gitlab-ci.yml</code> (http://docs.gitlab.com/ce/ci/yaml/README.html#gitlab-ci-yml)	File containing all definitions of how your project should be built
<code>script</code> (http://docs.gitlab.com/ce/ci/yaml/README.html#script)	Defines a shell script to be executed
<code>before_script</code> (http://docs.gitlab.com/ce/ci/yaml/README.html#before_script)	Used to define the command that should be run before (all) jobs

